

# Meet the DCAP Team

## Driving innovation

- Ayush
- Daksh
- Deepans
- h Pallavi

Together, our team, DCAP, is dedicated to developing and refining the Fine-Tuned Question-Answer Agent System, ensuring its robustness and efficiency.

# **Fine-Tuned Question-Answer Agent System**

Building an efficient and reliable AI for generating and solving Multiple-Choice Questions (MCQs) under strict constraints. This report details the architecture, methodology, and key innovations of our system.

# **Project Overview: Q-Agent & A-Agent**

This project developed a sophisticated Question-Answer (QA) agent system designed for specialized tasks:

- **Q-Agent:** Generates high-quality multiple-choice questions.
- **A-Agent:** Accurately solves those generated questions.

The system operates under stringent formatting, topic, and latency constraints, leveraging a self-play synthetic data generation pipeline and Supervised Fine-Tuning (SFT) with LoRA adapters for optimal efficiency and performance.

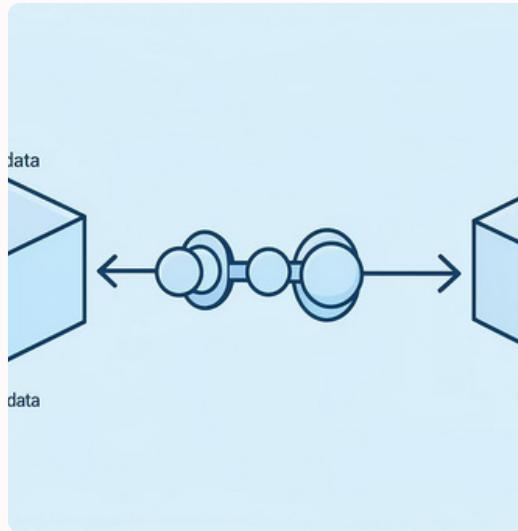
# Core Technologies: Foundation & Efficiency



**Qwen3-4B**

## Base Model: Qwen3-4B

Selected as the foundational model for both Q-Agent (question generation) and A-Agent (question answering). Its robust reasoning capabilities make it ideal for logical and puzzle-based domains, ensuring complex problem structures can be effectively processed.



## Fine-Tuning: Unslot (LoRA)

Utilized for parameter-efficient fine-tuning, significantly reducing memory usage and training time. The LoRA configuration (Rank=64, lora\_alpha=128, Dropout=0) targeted key modules to maintain model performance with minimal computational overhead.

# Training Methodology: SFT & Configuration

## Supervised Fine-Tuning (SFT)

Separate SFT was applied to optimize each agent:

- Q-Agent: Fine-tuned for precise MCQ generation.
- A-Agent: Fine-tuned for accurate MCQ solving.

Loss computation was strictly on assistant responses. The training configuration included:

- bf16 precision
- max\_seq\_length = 1024
- batch\_size = 32
- gradient\_accumulation = 2
- 2 epochs
- learning\_rate = 2e-4

# Data Engineering: Precision & Quality Control



## Python Pipeline

Robust preprocessing scripts.



## JSON Schema

Strict enforcement and auto-correction.



## Deduplication

Jaccard-style similarity for diversity.

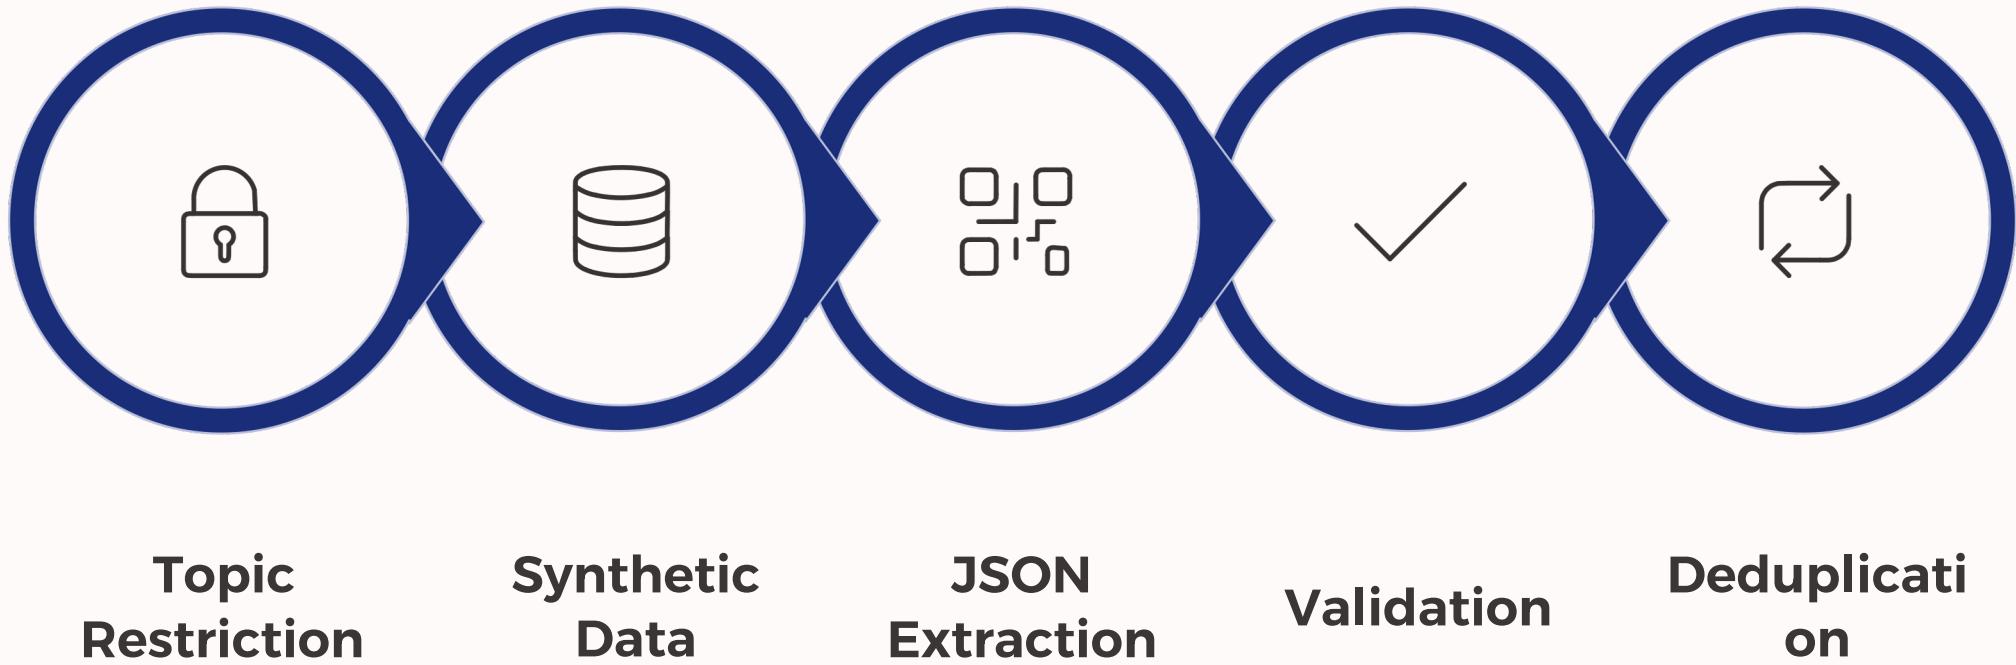


## Answer Balancing

Controlled A/B/C/D distribution.

Our comprehensive data engineering pipeline ensures high data quality, format integrity, and balanced answer distributions, critical for unbiased model training and optimal performance.

# Work Flow: From Topic to Fine-Tuning



The project followed a meticulous multi-step process, beginning with strict topic adherence and culminating in specialized fine-tuning for both agents.

# Adaptive Verification & Dataset Construction

## Adaptive Answer Verification

To ensure accuracy in logic-heavy domains (Seating, Blood Relations, Series), the teacher model solved questions twice, with majority voting. Syllogism verification was disabled due to low reliability, optimizing compute.

## Dataset Construction

Two specialized datasets were created:

- **A-Agent Dataset:** Input (Question + Choices), Output (JSON answer).
- **Q-Agent Dataset:** Input (Topic prompt), Output (JSON MCQ).

Both were saved in structured JSON format, ready for fine-tuning.

# Fine-Tuning and End-to-End Testing

1

## Separate LoRA Adapters

Dedicated LoRA adapters were retrained for each agent, allowing for specialized optimization without retraining the entire base model. This modular approach enhances efficiency and performance.

2

## Comprehensive Testing

The system underwent rigorous end-to-end testing, covering question generation, filter pass rates, answer prediction, and accuracy computation. Additional checks included conversation structure, JSON schema, and answer distribution validation.

# Key Design Strengths

- **Robust JSON Reliability Engineering**

Ensuring consistent and accurate JSON output.

- **Strict Schema Validation**

Guarding against malformed outputs and data inconsistencies.

- **Balanced Answer Distribution**

Preventing model bias toward specific answer choices.

- **Deduplication for Generalization**

Enhancing dataset diversity and model's ability to generalize.

- **Adaptive Correctness Verification**

Intelligently verifying answers for logical consistency.

- **Parameter-Efficient Fine-Tuning**

Optimizing computational resources without sacrificing performance.

- **Latency-Aware Inference Setup**

Ensuring quick response times for real-time applications.

# **Conclusion: A Robust QA System**

This project successfully developed a complete QA system pipeline, integrating synthetic dataset generation, stringent data validation, parameter-efficient fine-tuning, and robust reliability engineering. The result is a high-performing Question-Answer agent system that meets strict evaluation constraints and sets a new standard for efficiency and accuracy.